

AIM:- Data Loading, Data preprocessing, Data exploration, and Data preparation.

1. Data Loading:

- Load a dataset (Iris dataset: <https://www.kaggle.com/datasets/uciml/iris>) into your preferred ML environment (Python).

```
import pandas as pd
df = pd.read_csv("D:\\5th Sem\\LAB\\ML\\Iris.csv")
```

- Display the first few rows of the dataset to inspect its structure and content.

```
print("First 5 rows of the Iris dataset:-\n", df.head())
```

Output:-

```
First 5 rows of the Iris dataset:-
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0   1           5.1           3.5           1.4           0.2  Iris-setosa
1   2           4.9           3.0           1.4           0.2  Iris-setosa
2   3           4.7           3.2           1.3           0.2  Iris-setosa
3   4           4.6           3.1           1.5           0.2  Iris-setosa
4   5           5.0           3.6           1.4           0.2  Iris-setosa
```

- Check the dimensions of the dataset (number of rows and columns).

```
print("Dimension of the dataset: ", df.shape)
```

Output:-

```
Dimension of the dataset: (150, 6)
```

- Identify the data types of each column (numeric, categorical, text, etc.).

```
print("Data types of each column:\n", df.dtypes)
```

Output:-

```
Data types of each column:
   Id           int64
   SepalLengthCm  float64
   SepalWidthCm   float64
   PetalLengthCm  float64
   PetalWidthCm   float64
   Species        object
dtype: object
```

2. Data Exploration:

- Calculate basic summary statistics for the numeric columns (mean, median, min, max, standard deviation).

```
import matplotlib.pyplot as plt
import seaborn as sns
from DataLoading import df
numeric_columns = ['SepalLengthCm', 'SepalWidthCm',
                   'PetalLengthCm', 'PetalWidthCm']
print("Basic Statistics for the numeric columns:-\n",
      df[numeric_columns].describe())
```

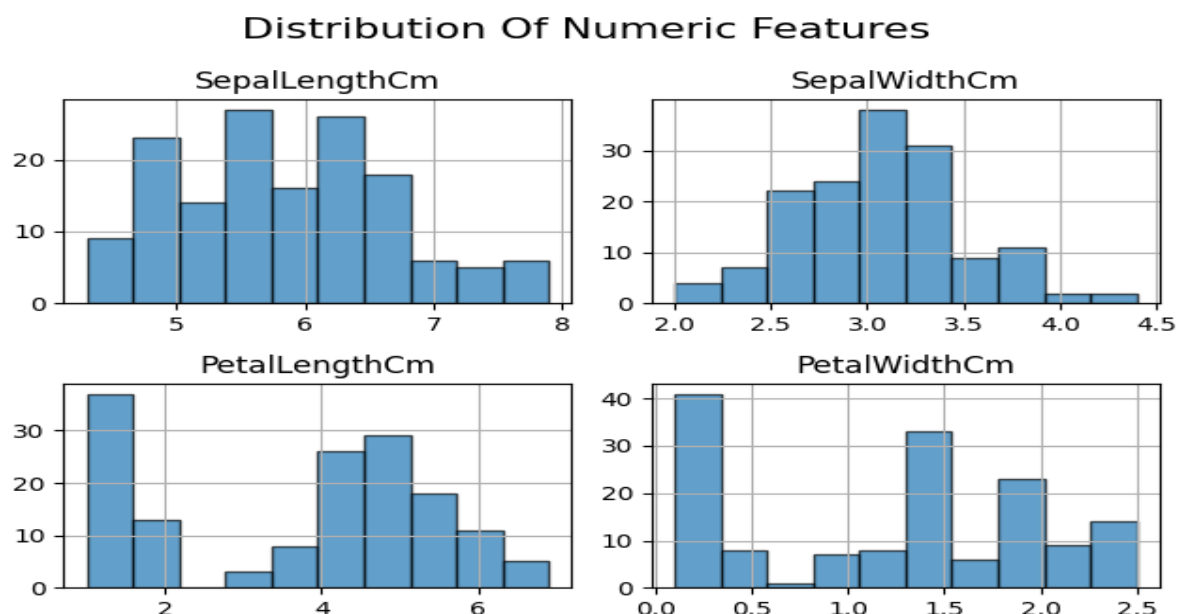
Output:-

```
Basic Statistics for the numeric columns:-
count      SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm
mean         5.843333      3.054000      3.758667      1.198667
std          0.828066      0.433594      1.764420      0.763161
min          4.300000      2.000000      1.000000      0.100000
25%          5.100000      2.800000      1.600000      0.300000
50%          5.800000      3.000000      4.350000      1.300000
75%          6.400000      3.300000      5.100000      1.800000
max          7.900000      4.400000      6.900000      2.500000
```

- Visualize the distribution of numeric features using histograms or box plots.

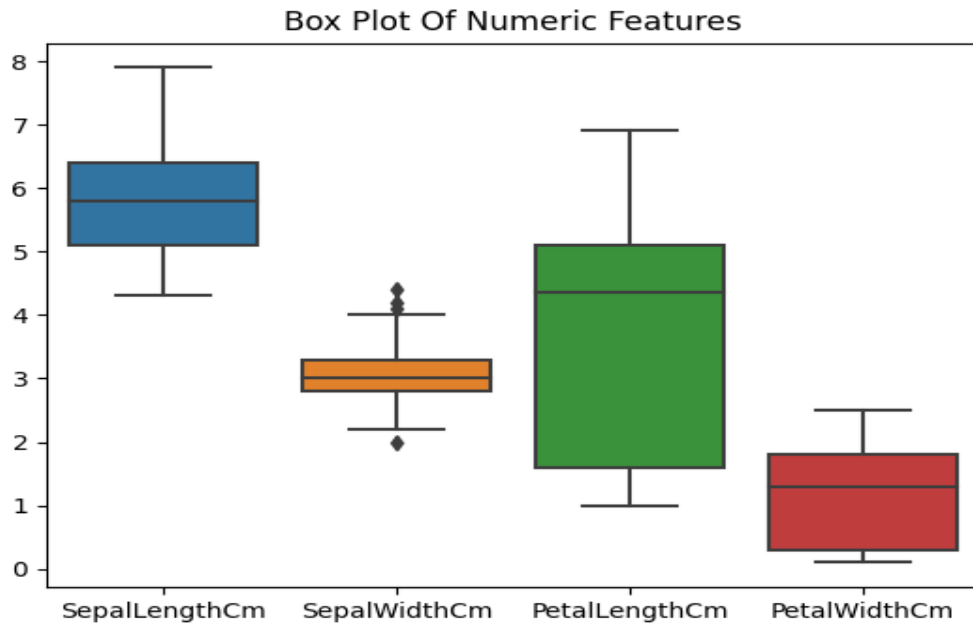
```
# Creating histograms for numeric features
df[numeric_columns].hist(edgecolor='black', alpha=0.7)
plt.suptitle("Distribution Of Numeric Features", fontsize=16)
plt.tight_layout()
plt.show()
```

Output:-



- # Creating box plots for numeric features
sns.boxplot(data=df[numeric_columns])
plt.title("Box Plot Of Numeric Features")
plt.show()

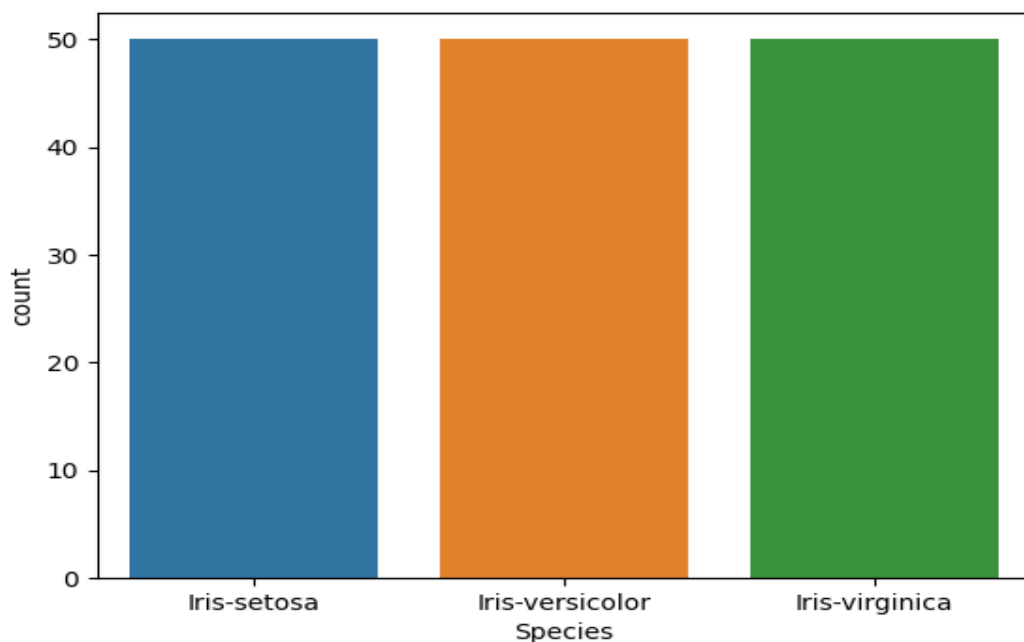
Output:-



- Explore the frequency distribution of categorical features using bar plots.

```
sns.countplot(data=df, x='Species')  
plt.show()
```

Output:-



3. Data Preprocessing:

- Handle missing values: Identify and handle any missing values in the dataset (e.g., imputation, removal).

```
from DataLoading import df
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
# Display the number of missing values in each column
missingValues = df.isnull().sum()
print("Missing values per column:-")
print(missingValues)
```

Output:-

```
Missing values per column:-
Id          0
SepalLengthCm  0
SepalWidthCm  0
PetalLengthCm  0
PetalWidthCm  0
Species      0
dtype: int64
```

- Display the number of missing values after imputation.

```
# Fill the missing values with the mean of each column except
column4
df_imputed = df.fillna(df.drop('Species', axis=1).mean())
missingValues_after_imputation = df_imputed.isnull().sum()
print("Missing Values after imputation:-")
print(missingValues_after_imputation)
```

Output:-

```
Missing Values after imputation:-
Id          0
SepalLengthCm  0
SepalWidthCm  0
PetalLengthCm  0
PetalWidthCm  0
Species      0
dtype: int64
```

- Encode categorical variables: Convert categorical variables into numerical form (e.g., one-hot encoding, label encoding).

```

# Categorical Column
categorical_column = 'Species'
# One-Hot Encoding (Creating Dummy Variables)
data_encoded_onehot = pd.get_dummies(df,
columns=[categorical_column], prefix=[categorical_column])
# Display the first few rows of the encoded data
print("One-Hot Encoded Data:")
print(data_encoded_onehot.head())
# Label Encoding
data_encoded_label = df.copy()
label_encoder = LabelEncoder()
data_encoded_label[categorical_column] =
label_encoder.fit_transform(data_encoded_label[categorical_column])
# Displays the first few rows of the encoded data
print("\nLabel Encoded Data:")
print(data_encoded_label.head())
# Reverse Label Encoding(for demonstration purposes)
reverse_encoded_labels =
label_encoder.inverse_transform(data_encoded_label[categorical_col
umn])
data_encoded_label[categorical_column] = reverse_encoded_labels
# Display the first few rows of the data with reversed level encoding
print("\nData with Reversed Label Encoding:")
print(data_encoded_label.head())

```

Output:-

```

One-Hot Encoded Data:
   Id  SepalLengthCm  ...  Species_Iris-versicolor  Species_Iris-virginica
0    1             5.1  ...                      False                      False
1    2             4.9  ...                      False                      False
2    3             4.7  ...                      False                      False
3    4             4.6  ...                      False                      False
4    5             5.0  ...                      False                      False

[5 rows x 8 columns]

Label Encoded Data:
   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species
0    1             5.1           3.5           1.4           0.2         0
1    2             4.9           3.0           1.4           0.2         0
2    3             4.7           3.2           1.3           0.2         0
3    4             4.6           3.1           1.5           0.2         0
4    5             5.0           3.6           1.4           0.2         0

```

Data with Reversed Label Encoding:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

- Feature scaling: Normalize or standardize numeric features to bring them to a similar scale.

```
numeric_columns = ['SepalLengthCm', 'SepalWidthCm',  
'PetalLengthCm', 'PetalWidthCm']  
# Standardization  
scaler_standard = StandardScaler()  
data_standardized =  
pd.DataFrame(scaler_standard.fit_transform(df[numeric_columns]),  
columns=numeric_columns)  
# Display the first few rows of standardized data  
print("Standardized Data:")  
print(data_standardized.head())  
  
# Normalization(MinMax Scaling)  
scaler_minmax = MinMaxScaler()  
data_normalized =  
pd.DataFrame(scaler_minmax.fit_transform(df[numeric_columns]),  
columns=numeric_columns)  
# Display the first few rows of normalized data  
print("\nNormalized Data:")  
print(data_normalized.head())
```

Output:-

```
Standardized Data:  
   SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  
0      -0.900681      1.032057     -1.341272     -1.312977  
1      -1.143017     -0.124958     -1.341272     -1.312977  
2      -1.385353      0.337848     -1.398138     -1.312977  
3      -1.506521      0.106445     -1.284407     -1.312977  
4      -1.021849      1.263460     -1.341272     -1.312977  
  
Normalized Data:  
   SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  
0      0.222222      0.625000      0.067797      0.041667  
1      0.166667      0.416667      0.067797      0.041667  
2      0.111111      0.500000      0.050847      0.041667  
3      0.083333      0.458333      0.084746      0.041667  
4      0.194444      0.666667      0.067797      0.041667  
  
Process finished with exit code 0
```

4. Data Preparation for ML:

- Split the dataset into training and testing sets (e.g., 80% for training, 20% for testing).

```
from DataLoading import df
from sklearn.model_selection import train_test_split
# Features and target variable
X = df.drop(columns=['Species']) # Features
Y = df['Species'] # Target Variable
# Split the dataset into training and testing sets(80% training, 20%testing)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
random_state=42)
# Display the shape of the training and testing sets
print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of Y_train:", Y_train.shape)
print("Shape of Y_test:", Y_test.shape)
```

Output:-

```
Shape of X_train: (120, 5)
Shape of X_test: (30, 5)
Shape of Y_train: (120,)
Shape of Y_test: (30,)
```

- Ensure the data is in the appropriate format for the ML algorithms (e.g., arrays, matrices).

```
# Convert the data to arrays or matrices
X_train_array = X_train.values
X_test_array = X_test.values
Y_train_array = Y_train.values
Y_test_array = Y_test.values
# Display the type and shape of the array
print("Type of X_train_array:", type(X_train_array))
print("Shape of X_train_array:", X_train_array.shape)
print("Type of Y_train_array:", type(Y_train_array))
print("Shape of Y_train_array:", X_train_array.shape)
```

Output:-

```
Type of X_train_array: <class 'numpy.ndarray'>
Shape of X_train_array: (120, 5)
Type of Y_train_array: <class 'numpy.ndarray'>
Shape of Y_train_array: (120, 5)
```

Submitted By,

Name- Shankar Singh Mahanty

Regd. No.- 2101020758

Roll No.- CSE21238

Group- 3

Sem- 5th

Branch- CSE